

3rd Meet

Long Time No See

□ 오늘 공부할 것은

02

사이킷런으로 시작하는 머신러닝

05. 데이터 전처리	116
데이터 인코딩	116
피처 스케일링과 정규화	122
StandardScaler	123
MinMaxScaler	125
학습 데이터와 테스트 데이터의 스케일링 변환 시 유의점	126
06. 사이킷런으로 수행하는 타이타닉 생존자 예측	129
07. 정리	143

Now, It's your turn !!!

40min ~ 50min

[02-05] 데이터 전처리 (Data Preprocessing)

- 데이터 인코딩 (Encoding)
 - . Label Encoding
 - . One-Hot Encoding
- 피처 스케일링과 정규화 (Standardization and Normalization)
 - . StandardScaler
 - . MinMaxScaler
 - . 학습 데이터와 테스트 데이터의 스케일링 변환 시 유의점

[02-06] 사이킷런으로 수행하는 타이타닉 생존자 예측

[02-07] 정리

[02-05]

데이터 전처리

(Data Preprocessing)

□ Data Preprocessing

- 결손값 / 결측치 / NaN / Null → 버리거나 평균치 등을 이용한 대체
- 문자열 / String → 숫자형으로 변환(encoding)
 - . 카테고리형 → 코드 값으로 대체
 - . 텍스트형 → 피처 벡터화 (feature vectorization) 또는 삭제

□ Encoding - Label encoding

- Category → 코드형 숫자 값으로 변환
- But 숫자의 크기는?
 - . 트리 구조에서는 별 문제가 없으나, 선형 회귀에서는 문제

LabelEncoder

```
from sklearn.preprocessing import LabelEncoder

items = ['TV', '냉장고', '전자렌지', '컴퓨터', '선풍기', '선풍기', '믹서', '믹서']

encoder = LabelEncoder()
labels = encoder.fit_transform(items)

print('인코딩 변환값:', labels)
print('인코딩 클래스:', encoder.classes_)
```

```
인코딩 변환값: [0 1 4 5 3 3 2 2]
인코딩 클래스: ['TV' '냉장고' '믹서' '선풍기' '전자렌지' '컴퓨터']
```

```
original = encoder.inverse_transform([4, 5, 2, 0, 1, 1, 3, 3])
```

```
print('디코딩 원본 값:', original)
```

```
디코딩 원본 값: ['전자렌지' '컴퓨터' '믹서' 'TV' '냉장고' '냉장고' '선풍기' '선풍기']
```


□ Encoding - One-Hot Encoding

- 입력 = numpy 행렬

```
from sklearn.preprocessing import OneHotEncoder
import numpy as np

items = np.array(['TV', '냉장고', '전자렌지', '컴퓨터', '선종기', '선종기', '믹서', '믹서'])
items

array(['TV', '냉장고', '전자렌지', '컴퓨터', '선종기', '선종기', '믹서', '믹서'], dtype='<U4')

items = items.reshape(-1,1)
items

array([[ 'TV'],
       [ '냉장고'],
       [ '전자렌지'],
       [ '컴퓨터'],
       [ '선종기'],
       [ '선종기'],
       [ '믹서'],
       [ '믹서']], dtype='<U4')
```

```
oh_encoder = OneHotEncoder()
oh_labels = oh_encoder.fit_transform(items)

print('원-핫 인코딩 데이터')
print(oh_labels.toarray())

print('원-핫 인코딩 데이터 차원')
print(oh_labels.shape)

원-핫 인코딩 데이터
[[1.  0.  0.  0.  0.  0.]
 [0.  1.  0.  0.  0.  0.]
 [0.  0.  0.  0.  1.  0.]
 [0.  0.  0.  0.  0.  1.]
 [0.  0.  0.  1.  0.  0.]
 [0.  0.  0.  1.  0.  0.]
 [0.  0.  1.  0.  0.  0.]
 [0.  0.  1.  0.  0.  0.]]
원-핫 인코딩 데이터 차원
(8, 6)

original = oh_encoder.inverse_transform([[0, 0, 0, 0, 1, 0], [0, 0, 0, 0, 0, 1]])

print('디코딩 원본 값:', original)

디코딩 원본 값: [['전자렌지']
                  ['컴퓨터']]
```

□ Encoding - One-Hot Encoding

```
import pandas as pd

df = pd.DataFrame({'item': ['TV', '냉장고', '전자렌지', '컴퓨터', '선풍기', '선풍기', '믹서', '믹서'] })
df
```

	item
0	TV
1	냉장고
2	전자렌지
3	컴퓨터
4	선풍기
5	선풍기
6	믹서
7	믹서

```
oh_labels = pd.get_dummies(df)
oh_labels
```

	item_TV	item_냉장고	item_믹서	item_선풍기	item_전자렌지	item_컴퓨터
0	True	False	False	False	False	False
1	False	True	False	False	False	False
2	False	False	False	False	True	False
3	False	False	False	False	False	True
4	False	False	False	True	False	False
5	False	False	False	True	False	False
6	False	False	True	False	False	False
7	False	False	True	False	False	False

□ Feature Scaling - Standardization (표준화)

- 평균 = 0, 분산 = 1, Gaussian 분포

$$x_{i_new} = \frac{x_i - mean(x)}{stdev(x)}$$

- Gaussian 분포 기반 (in 사이킷런)
 - . SVM (Support Vector Machine)
 - . Linear Regression
 - . Logistic Regression

□ Feature Scaling - Standardization (표준화) - StandardScaler

```
from sklearn.datasets import load_iris
import pandas as pd

iris = load_iris()
iris_data = iris.data
iris_df = pd.DataFrame(data=iris_data, columns=iris.feature_names)

print('feature 들의 분산 값')
print(iris_df.var())
```

```
feature 들의 분산 값
sepal length (cm)    0.685694
sepal width (cm)     0.189979
petal length (cm)    3.116278
petal width (cm)     0.581006
dtype: float64
```

```
iris_df.describe()
```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)
count	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.057333	3.758000	1.199333
std	0.828066	0.435866	1.765298	0.762238
min	4.300000	2.000000	1.000000	0.100000
25%	5.100000	2.800000	1.600000	0.300000
50%	5.800000	3.000000	4.350000	1.300000
75%	6.400000	3.300000	5.100000	1.800000
max	7.900000	4.400000	6.900000	2.500000

```
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
iris_std = scaler.fit_transform(iris_df)

iris_df_std = pd.DataFrame(data=iris_std, columns=iris.feature_names)

print('feature 들의 분산 값')
print(iris_df_std.var())
```

```
feature 들의 분산 값
sepal length (cm)    1.006711
sepal width (cm)     1.006711
petal length (cm)    1.006711
petal width (cm)     1.006711
dtype: float64
```

```
iris_df_std.describe()
```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)
count	1.500000e+02	1.500000e+02	1.500000e+02	1.500000e+02
mean	-1.468455e-15	-1.823726e-15	-1.610564e-15	-9.473903e-16
std	1.003350e+00	1.003350e+00	1.003350e+00	1.003350e+00
min	-1.870024e+00	-2.433947e+00	-1.567576e+00	-1.447076e+00
25%	-9.006812e-01	-5.923730e-01	-1.226552e+00	-1.183812e+00
50%	-5.250608e-02	-1.319795e-01	3.364776e-01	1.325097e-01
75%	6.745011e-01	5.586108e-01	7.627583e-01	7.906707e-01
max	2.492019e+00	3.090775e+00	1.785832e+00	1.712096e+00

□ Feature Scaling - Normalization (정규화)

- 서로 다른 피처의 크기를 통일하기 위해 크기 변환

$$x_{i_new} = \frac{x_i - \min(x)}{\max(x) - \min(x)}$$

- '사이킷런'에서는 선형대수의 Normalization 방식 (= Vector Normalization)

$$x_{i_new} = \frac{x_i}{\sqrt{x_i^2 + y_i^2 + z_i^2}}$$

□ Feature Scaling - Normalization (정규화) - MinMaxScaler

- 데이터의 값을 0과 1 사이의 범위로 변환 (음수가 있다면 -1에서 1 사이)

```
from sklearn.datasets import load_iris
import pandas as pd

iris = load_iris()
iris_data = iris.data
iris_df = pd.DataFrame(data=iris_data, columns=iris.feature_names)

print('feature 들의 분산 값')
print(iris_df.var())
```

```
feature 들의 분산 값
sepal length (cm)    0.685694
sepal width (cm)     0.189979
petal length (cm)    3.116278
petal width (cm)     0.581006
dtype: float64
```

```
iris_df.describe()
```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)
count	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.057333	3.758000	1.199333
std	0.828066	0.435866	1.765298	0.762238
min	4.300000	2.000000	1.000000	0.100000
25%	5.100000	2.800000	1.600000	0.300000
50%	5.800000	3.000000	4.350000	1.300000
75%	6.400000	3.300000	5.100000	1.800000
max	7.900000	4.400000	6.900000	2.500000

```
from sklearn.preprocessing import MinMaxScaler
```

```
scaler = MinMaxScaler()
iris_mm = scaler.fit_transform(iris_df)
```

```
iris_df_mm = pd.DataFrame(data=iris_mm, columns=iris.feature_names)
```

```
print('feature들의 최대 값')
print(iris_df_mm.max())
```

```
feature들의 최대 값
sepal length (cm)    1.0
sepal width (cm)     1.0
petal length (cm)    1.0
petal width (cm)     1.0
dtype: float64
```

```
iris_df_mm.describe()
```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)
count	150.000000	150.000000	150.000000	150.000000
mean	0.428704	0.440556	0.467458	0.458056
std	0.230018	0.181611	0.299203	0.317599
min	0.000000	0.000000	0.000000	0.000000
25%	0.222222	0.333333	0.101695	0.083333
50%	0.416667	0.416667	0.567797	0.500000
75%	0.583333	0.541667	0.694915	0.708333
max	1.000000	1.000000	1.000000	1.000000

□ fit() + transform()

```
from sklearn.preprocessing import MinMaxScaler
import numpy as np

train_array = np.arange(0, 11).reshape(-1, 1)
train_array

array([[ 0],
       [ 1],
       [ 2],
       [ 3],
       [ 4],
       [ 5],
       [ 6],
       [ 7],
       [ 8],
       [ 9],
       [10]])

scaler = MinMaxScaler()

scaler.fit(train_array)

train_scaled = scaler.transform(train_array)
train_scaled

array([[0. ],
       [0.1],
       [0.2],
       [0.3],
       [0.4],
       [0.5],
       [0.6],
       [0.7],
       [0.8],
       [0.9],
       [1. ]])
```

```
test_array = np.arange(0, 6).reshape(-1, 1)
test_array

array([[0],
       [1],
       [2],
       [3],
       [4],
       [5]])

scaler.fit(test_array)

test_scaled = scaler.transform(test_array)
test_scaled

array([[0. ],
       [0.2],
       [0.4],
       [0.6],
       [0.8],
       [1. ]])
```

```
scaler = MinMaxScaler()

scaler.fit(train_array)

train_scaled = scaler.transform(train_array)
train_scaled

array([[0. ],
       [0.1],
       [0.2],
       [0.3],
       [0.4],
       [0.5],
       [0.6],
       [0.7],
       [0.8],
       [0.9],
       [1. ]])

test_scaled = scaler.transform(test_array)
test_scaled

array([[0. ],
       [0.1],
       [0.2],
       [0.3],
       [0.4],
       [0.5]])
```


- 같은 fit()을 이용한 transform() 중요 !!!

[02-06]

사이킷런으로 수행하는 타이타닉 생존자 예측

RMS Titanic

- if the tītanīc sank today...



https://www.reddit.com/r/Duklock/comments/bexm2x/if_the_titanic_sank_today/

- 1912년 4월 14일 밤 침몰. 총 2224명 중 710명 구조, 1514명 사망



https://ko.wikipedia.org/wiki/RMS_타이타닉#/media/파일:Titanic_lifeboat.jpg

https://ko.wikipedia.org/wiki/RMS_타이타닉#/media/파일:Titanic-New_York_Herald_front_page.jpeg



Information

- 탑승자 및 배에 대한 다양한 정보를 얻을 수 있다.

WHITE STAR LINE
ROYAL AND UNITED STATES MAIL STEAMERS.

Male Berth.....
Female Berth.....
Married Berth.....

ISMAI, IMRIE & CO.,
1 COCKSPUR STREET, S.W.,
25 LEADENHALL STREET, E.C.,
LONDON,
30, JAMES STREET
LIVERPOOL,
AND
CANUTE ROAD, SOUTHAMPTON.

Agent at PARIS—
NICHOLAS MARTIN, 9, Rue Scribe.

WHITE STAR LINE.
19, VIA ALLA NUZIATA GENOA
21, PIAZZA DELLA BORSA NAPLES
24, ST. MARK STREET BOSTON
9, BROADWAY NEW YORK
63, D'ALHOUSIE STREET QUEBEC
BELL TELEPHONE BUILDING
119, NOTRE DAME STREET WEST } MONTREAL

JAMES SCOTT & CO., Agents.
QUEENSTOWN.

OCEANIC STEAM NAVIGATION COMPANY, LIMITED, OF GREAT BRITAIN.
THIRD CLASS (Steerage) PASSENGER'S CONTRACT TICKET.
(NOT TRANSFERABLE.)

SHIP R.M.S. TITANIC of 45,000 Tons Register,
to take in Passengers at QUEENSTOWN for NEW YORK
on the 11 day of APRIL 19 1912

NAMES.	AGE.	No. of State Adults
<i>Patrick Ryan</i>	32	3
<i>W. J. Collins</i>		

Deposit.....£
Balance.....£
Total.....£

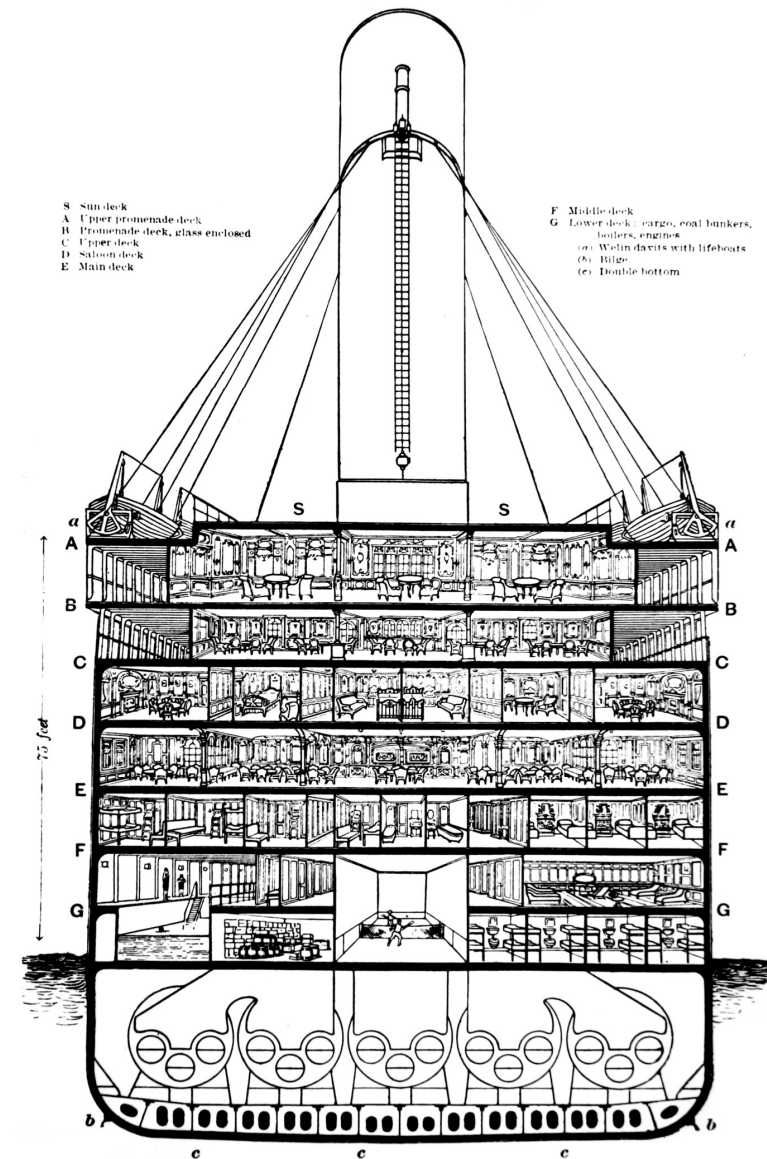
(To be paid at the Port
of Embarkation.)

BREAKFAST AT EIGHT O'CLOCK.—Oatmeal Porridge and Milk, Tea, Coffee, Sugar, Milk, Fresh Bread and Butter, HERRINGS, Potatoes, Ling Fish and Egg Sauce, or Irish Stew, according to the day of the week.
DINNER AT ONE O'CLOCK.—Soup, Beef, Mutton, Carrots and Turnips, Green Peas, Pork or Ling Fish and Sauce with Bread and Potatoes, Plum Pudding, Rice Pudding, Stewed Apples and Rice, or Stewed Prunes and Rice, according to the day of the week.
TEA AT SIX O'CLOCK.—Tea, Sugar, Milk, Fresh Bread and Butter, Jam, or Cheese and Pickles, Corned Beef, Hash, or Tinned Beef, according to the day of the week. Oatmeal Gruel will be supplied at 8 p.m.
NOTE.—For Women and Children—Chicken Broth or Beef Tea Daily at 11 a.m., and Tea, Coffee and Milk at all hours.

Mess Utensils and Bedding provided by the Ship.

For and on behalf of the OCEANIC STEAM NAVIGATION COMPANY, LIMITED, OF GREAT BRITAIN.

NOTICE TO STEERAGE PASSENGERS.—1.—If Steerage Passengers, through no default of their own, are not received on board on the day named in their Contract Tickets, or fail to obtain a Passage in the Ship, they should apply to the Migration Officer at the port, who will assist in obtaining redress, under the Merchant Shipping Acts.—2.—Third Class (Steerage) Passengers should carefully keep this part of their Contract Ticket till the end of the Voyage. This Contract Ticket is exempt from Stamp Duty.



https://twitter.com/weird_hist/status/912449385346281472

https://upload.wikimedia.org/wikipedia/commons/8/84/Titanic_cutaway_diagram.png

Problem

- Predict survival on the Titanic and get familiar with ML basics

Data Dictionary

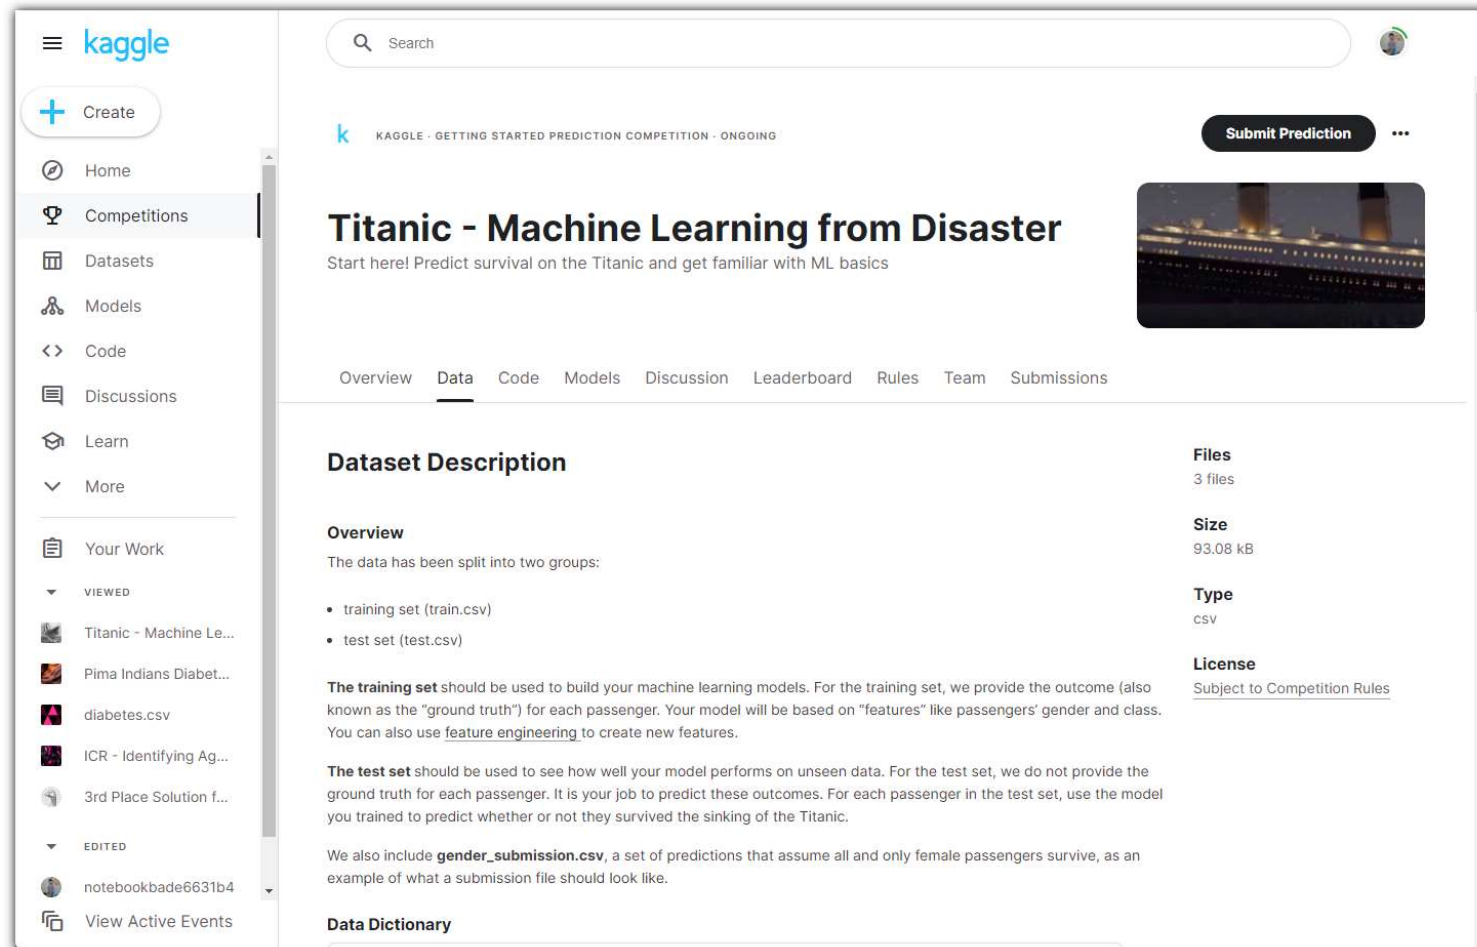
Variable	Definition	Key
survival	Survival	0 = No, 1 = Yes
pclass	Ticket class	1 = 1st, 2 = 2nd, 3 = 3rd
sex	Sex	
Age	Age in years	
sibsp	# of siblings / spouses aboard the Titanic	
parch	# of parents / children aboard the Titanic	
ticket	Ticket number	
fare	Passenger fare	
cabin	Cabin number	
embarked	Port of Embarkation	C = Cherbourg, Q = Queenstown, S = Southampton



<https://www.kaggle.com/c/titanic/data>

https://en.wikipedia.org/wiki/Rules_of_Survival#/media/File:Rules_of_Survival_Google_Play_Logo.png

□ Kaggle



The screenshot shows the Kaggle website interface for the 'Titanic - Machine Learning from Disaster' competition. The left sidebar contains navigation links: Home, Competitions (selected), Datasets, Models, Code, Discussions, Learn, and More. Below these are 'Your Work' and 'View Active Events' sections. The main content area has a search bar and a 'Submit Prediction' button. The competition title 'Titanic - Machine Learning from Disaster' is prominently displayed, followed by the subtitle 'Start here! Predict survival on the Titanic and get familiar with ML basics'. A navigation bar below the title includes links for Overview, Data (selected), Code, Models, Discussion, Leaderboard, Rules, Team, and Submissions. The 'Dataset Description' section provides an overview of the data, which is split into training and test sets. It lists the files: train.csv and test.csv. The description also mentions the 'gender_submission.csv' file as an example of a submission. On the right side of the dataset description, there are summary statistics: 3 files, 93.08 kB, and CSV type. The license is noted as 'Subject to Competition Rules'.

Titanic - Machine Learning from Disaster
Start here! Predict survival on the Titanic and get familiar with ML basics

Overview **Data** Code Models Discussion Leaderboard Rules Team Submissions

Dataset Description

Files
3 files

Size
93.08 kB

Type
csv

License
[Subject to Competition Rules](#)

Overview
The data has been split into two groups:

- training set (train.csv)
- test set (test.csv)

The training set should be used to build your machine learning models. For the training set, we provide the outcome (also known as the "ground truth") for each passenger. Your model will be based on "features" like passengers' gender and class. You can also use [feature engineering](#) to create new features.

The test set should be used to see how well your model performs on unseen data. For the test set, we do not provide the ground truth for each passenger. It is your job to predict these outcomes. For each passenger in the test set, use the model you trained to predict whether or not they survived the sinking of the Titanic.

We also include **gender_submission.csv**, a set of predictions that assume all and only female passengers survive, as an example of what a submission file should look like.

Data Dictionary

※ 출처 : <https://www.kaggle.com/competitions/titanic/data>

□ 기본 정보 확인

```
import pandas as pd
```

```
df = pd.read_csv('./titanic_train.csv')  
df.head()
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN	S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	C85	C
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C123	S
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	NaN	S

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 891 entries, 0 to 890  
Data columns (total 12 columns):  
#   Column      Non-Null Count  Dtype  
---  ---  
0   PassengerId  891 non-null    int64  
1   Survived     891 non-null    int64  
2   Pclass       891 non-null    int64  
3   Name         891 non-null    object  
4   Sex          891 non-null    object  
5   Age          714 non-null    float64  
6   SibSp        891 non-null    int64  
7   Parch        891 non-null    int64  
8   Ticket       891 non-null    object  
9   Fare         891 non-null    float64  
10  Cabin        204 non-null    object  
11  Embarked     889 non-null    object  
dtypes: float64(2), int64(5), object(5)  
memory usage: 83.7+ KB
```

□ 결측치 처리


```
df.isnull().sum()
```

```
PassengerId    0
Survived        0
Pclass          0
Name            0
Sex             0
Age            177
SibSp           0
Parch           0
Ticket          0
Fare            0
Cabin          687
Embarked        2
dtype: int64
```

```
import matplotlib.pyplot as plt
```

```
plt.figure(figsize=(8,3))
df.isnull().sum().plot.barh()
```

<Axes: >



```
df['Age'].fillna(df['Age'].mean(), inplace=True)
df['Cabin'].fillna('N', inplace=True)
df['Embarked'].fillna('N', inplace=True)
```

```
df.isnull().sum().sum()
```

```
0
```

□ 문자열 처리

```
df['Sex'].value_counts()
```

```
Sex
male    577
female  314
Name: count, dtype: int64
```

```
df['Embarked'].value_counts()
```

```
Embarked
S    644
C    168
Q     77
Name: count, dtype: int64
```

```
df['Cabin'].value_counts()
```

```
Cabin
B96 B98    4
G6         4
C23 C25 C27 4
C22 C26    3
F33        3
..
E34        1
C7         1
C54        1
E36        1
C148       1
Name: count, Length: 147, dtype: int64
```

```
df['Cabin'] = df['Cabin'].str[:1]
df['Cabin'].value_counts()
```

```
Cabin
C    59
B    47
D    33
E    32
A    15
F    13
G     4
T     1
Name: count, dtype: int64
```

```
from sklearn import preprocessing

def encode_features(dataDF):
    features = ['Cabin', 'Sex', 'Embarked']
    for feature in features:
        le = preprocessing.LabelEncoder()
        le = le.fit(dataDF[feature])
        dataDF[feature] = le.transform(dataDF[feature])

    return dataDF

df = encode_features(df)
df.head()
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr. Owen Harris	1	22.0	1	0	A/5 21171	7.2500	7	3
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	0	38.0	1	0	PC 17599	71.2833	2	0
2	3	1	3	Heikkinen, Miss. Laina	0	26.0	0	0	STON/O2. 3101282	7.9250	7	3
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	0	35.0	1	0	113803	53.1000	2	3
4	5	0	3	Allen, Mr. William Henry	1	35.0	0	0	373450	8.0500	7	3

□ 데이터 탐색



```
df.groupby(['Sex', 'Survived'])['Survived'].count()

Sex    Survived
female 0         81
       1        233
male   0        468
       1        109
Name: Survived, dtype: int64

import seaborn as sns

fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(10,3))
sns.barplot(x='Sex', y='Survived', data=df, ax=axes[0])
sns.barplot(x='Pclass', y='Survived', hue='Sex', data=df, ax=axes[1])


<Axes: xlabel='Pclass', ylabel='Survived'>
```



선실 등급에 따른 생존율
여성의 생존율

□ 카테고리화

```
def get_category(age):  
    cat = ''  
    if age <= -1: cat = 'Unknown'  
    elif age <= 5: cat = 'Baby'  
    elif age <= 12: cat = 'Child'  
    elif age <= 18: cat = 'Teenager'  
    elif age <= 25: cat = 'Student'  
    elif age <= 35: cat = 'Young Adult'  
    elif age <= 60: cat = 'Adult'  
    else : cat = 'Elderly'  
  
    return cat  
  
df['Age_cat'] = df['Age'].apply(lambda x : get_category(x))  
  
plt.figure(figsize=(8,5))  
group_names = ['Unknown', 'Baby', 'Child', 'Teenager', 'Student', 'Young Adult', 'Adult', 'Elderly']  
sns.barplot(x='Age_cat', y = 'Survived', hue='Sex', data=df, order=group_names)  
  
df.drop('Age_cat', axis=1, inplace=True)
```



어린 아이들의 생존율

여성의 생존율

□ Training

```
from sklearn.model_selection import train_test_split

y_df = df['Survived']
X_df = df.drop(['Survived', 'PassengerId', 'Name', 'Ticket'], axis=1)

X_train, X_test, y_train, y_test = train_test_split(X_df, y_df, test_size=0.2, random_state=11)
```

DecisionTreeClassifier

```
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score

dt_clf = DecisionTreeClassifier(random_state=11)

dt_clf.fit(X_train, y_train)
dt_pred = dt_clf.predict(X_test)

accuracy_score(y_test, dt_pred)

0.7877094972067039
```

RandomForestClassifier

```
from sklearn.ensemble import RandomForestClassifier

rf_clf = RandomForestClassifier(random_state=11)

rf_clf.fit(X_train, y_train)
rf_pred = rf_clf.predict(X_test)

accuracy_score(y_test, rf_pred)

0.8547486033519553
```

LogisticRegression

```
from sklearn.linear_model import LogisticRegression

lr_clf = LogisticRegression(max_iter=200)

lr_clf.fit(X_train, y_train)
lr_pred = lr_clf.predict(X_test)

accuracy_score(y_test, lr_pred)

0.8491620111731844
```

KFold

```
import numpy as np
from sklearn.model_selection import KFold

def exec_kfold(clf, folds=5):
    kfold = KFold(n_splits=folds)
    scores = []

    for iter_count, (train_index, test_index) in enumerate(kfold.split(X_df)):
        X_train, X_test = X_df.values[train_index], X_df.values[test_index]
        y_train, y_test = y_df.values[train_index], y_df.values[test_index]

        clf.fit(X_train, y_train)
        predictions = clf.predict(X_test)
        accuracy = accuracy_score(y_test, predictions)
        scores.append(accuracy)
        print("교차 검증 {0} 정확도: {1:.4f}".format(iter_count, accuracy))

    return np.mean(scores)
```

```
exec_kfold(dt_clf, folds=5)
```

```
교차 검증 0 정확도: 0.7542
교차 검증 1 정확도: 0.7809
교차 검증 2 정확도: 0.7865
교차 검증 3 정확도: 0.7697
교차 검증 4 정확도: 0.8202
0.782298662984119
```

CrossValidation

```
from sklearn.model_selection import cross_val_score

scores = cross_val_score(dt_clf, X_df, y_df, cv=5)
for iter_count, accuracy in enumerate(scores):
    print("교차 검증 {0} 정확도: {1:.4f}".format(iter_count, accuracy))

np.mean(scores)
```

```
교차 검증 0 정확도: 0.7430
교차 검증 1 정확도: 0.7753
교차 검증 2 정확도: 0.7921
교차 검증 3 정확도: 0.7865
교차 검증 4 정확도: 0.8427
0.7879291946519366
```

GridSearchCV

```
from sklearn.model_selection import GridSearchCV

parameters = {'max_depth':[2,3,5,10],
              'min_samples_split':[2,3,5],
              'min_samples_leaf':[1,5,8]}

grid_dclf = GridSearchCV(dt_clf, param_grid=parameters, scoring='accuracy', cv=5)
grid_dclf.fit(X_train, y_train)

print('GridSearchCV 최적 하이퍼 파라미터:', grid_dclf.best_params_)
print('GridSearchCV 최고 정확도: {0:.4f}'.format(grid_dclf.best_score_))
best_dclf = grid_dclf.best_estimator_

dpredictions = best_dclf.predict(X_test)
accuracy = accuracy_score(y_test, dpredictions)
accuracy
```

```
GridSearchCV 최적 하이퍼 파라미터: {'max_depth': 3, 'min_samples_leaf': 5, 'min_samples_split': 2}
GridSearchCV 최고 정확도: 0.7992
0.8715083798882681
```

You've really worked hard today

Next Week ~ ?